

## a2/kmeans/omp\_naive\_kmeans.c

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include "kmeans.h"
4 /*
5  * TODO: include openmp header file
6  */
7 #include <omp.h>
8
9 // square of Euclid distance between two multi-dimensional points
10 inline static double euclid_dist_2(int numdims, /* no. dimensions */
11                                 double *coord1, /* [numdims] */
12                                 double *coord2) /* [numdims] */
13 {
14     int i;
15     double ans = 0.0;
16
17     for (i = 0; i < numdims; i++)
18         ans += (coord1[i] - coord2[i]) * (coord1[i] - coord2[i]);
19
20     return ans;
21 }
22
23 inline static int find_nearest_cluster(int numClusters, /* no. clusters */
24                                         int numCoords, /* no. coordinates */
25                                         double *object, /* [numCoords] */
26                                         double *clusters) /* [numClusters][numCoords] */
27 {
28     int index, i;
29     double dist, min_dist;
30
31     // find the cluster id that has min distance to object
32     index = 0;
33     min_dist = euclid_dist_2(numCoords, object, clusters);
34
35     for (i = 1; i < numClusters; i++)
36     {
37         dist = euclid_dist_2(numCoords, object, &clusters[i * numCoords]);
38         // no need square root
39         if (dist < min_dist)
40             { // find the min and its array index
41                 min_dist = dist;
42                 index = i;
43             }
44     }
45     return index;
46 }
47
48 void kmeans(double *objects, /* in: [numObjs][numCoords] */
49             int numCoords, /* no. coordinates */
50             int numObjs, /* no. objects */
51             int numClusters, /* no. clusters */

```

```

52     double threshold,      /* minimum fraction of objects that change membership */
53     long loop_threshold, /* maximum number of iterations */
54     int *membership,      /* out: [numObjs] */
55     double *clusters)     /* out: [numClusters][numCoords] */
56 {
57     int i, j;
58     int index, loop = 0;
59     double timing = 0;
60
61     double delta;          // fraction of objects whose clusters change in each loop
62     int *newClusterSize; // [numClusters]: no. objects assigned in each new cluster
63     double *newClusters; // [numClusters][numCoords]
64     int nthreads;         // no. threads
65
66     nthreads = omp_get_max_threads();
67     printf("OpenMP Kmeans - Naive\t(number of threads: %d)\n", nthreads);
68
69     // initialize membership
70     for (i = 0; i < numObjs; i++)
71         membership[i] = -1;
72
73     // initialize newClusterSize and newClusters to all 0
74     newClusterSize = (typeof(newClusterSize))calloc(numClusters, sizeof(*newClusterSize));
75     newClusters = (typeof(newClusters))calloc(numClusters * numCoords,
76         sizeof(*newClusters));
76
77     timing = wtime();
78
79     do
80     {
81         // before each loop, set cluster data to 0
82         for (i = 0; i < numClusters; i++)
83         {
84             for (j = 0; j < numCoords; j++)
85                 newClusters[i * numCoords + j] = 0.0;
86             newClusterSize[i] = 0;
87         }
88
89         delta = 0.0;
90
91     /*
92     * TODO: Detect parallelizable region and use appropriate OpenMP pragmas
93     */
94     #pragma omp parallel for private(index, j)
95         for (i = 0; i < numObjs; i++)
96         {
97             // find the array index of nearest cluster center
98             index = find_nearest_cluster(numClusters, numCoords, &objects[i * numCoords],
99             clusters);
100
101             // if membership changes, increase delta by 1
102             if (membership[i] != index)
103             {

```

```
104 #pragma omp atomic // protect update on shared "delta" variable
105         delta += 1.0;
106     }
107
108     // assign the membership to object i
109     membership[i] = index;
110
111 // update new cluster centers : sum of objects located within
112 /*
113 * TODO: protect update on shared "newClusterSize" array
114 */
115 #pragma omp atomic
116         newClusterSize[index]++;
117         for (j = 0; j < numCoords; j++)
118     /*
119     * TODO: protect update on shared "newClusters" array
120     */
121 #pragma omp atomic
122         newClusters[index * numCoords + j] += objects[i * numCoords + j];
123     }
124
125     // average the sum and replace old cluster centers with newClusters
126     for (i = 0; i < numClusters; i++)
127     {
128         if (newClusterSize[i] > 0)
129         {
130             for (j = 0; j < numCoords; j++)
131             {
132                 clusters[i * numCoords + j] = newClusters[i * numCoords + j] /
133             newClusterSize[i];
134             }
135         }
136     }
137
138     // Get fraction of objects whose membership changed during this loop. This is used
139     // as a convergence criterion.
140     delta /= numObjs;
141
142     loop++;
143     printf("\r\tcompleted loop %d", loop);
144     fflush(stdout);
145 } while (delta > threshold && loop < loop_threshold);
146     timing = wtime() - timing;
147     printf("\n          nloops = %3d    (total = %7.4fs)  (per loop = %7.4fs)\n", loop,
148 timing, timing / loop);
149
150     free(newClusters);
151     free(newClusterSize);
152 }
```