

## a2/kmeans/omp\_reduction\_kmeans.c

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include "kmeans.h"
4 /*
5  * TODO: include openmp header file
6  */
7 #include <omp.h>
8
9 // square of Euclid distance between two multi-dimensional points
10 inline static double euclid_dist_2(int numdims, /* no. dimensions */
11                                 double *coord1, /* [numdims] */
12                                 double *coord2) /* [numdims] */
13 {
14     int i;
15     double ans = 0.0;
16
17     for (i = 0; i < numdims; i++)
18         ans += (coord1[i] - coord2[i]) * (coord1[i] - coord2[i]);
19
20     return ans;
21 }
22
23 inline static int find_nearest_cluster(int numClusters, /* no. clusters */
24                                         int numCoords, /* no. coordinates */
25                                         double *object, /* [numCoords] */
26                                         double *clusters) /* [numClusters][numCoords] */
27 {
28     int index, i;
29     double dist, min_dist;
30
31     // find the cluster id that has min distance to object
32     index = 0;
33     min_dist = euclid_dist_2(numCoords, object, clusters);
34
35     for (i = 1; i < numClusters; i++)
36     {
37         dist = euclid_dist_2(numCoords, object, &clusters[i * numCoords]);
38         // no need square root
39         if (dist < min_dist)
40             { // find the min and its array index
41                 min_dist = dist;
42                 index = i;
43             }
44     }
45     return index;
46 }
47
48 void kmeans(double *objects, /* in: [numObjs][numCoords] */
49             int numCoords, /* no. coordinates */
50             int numObjs, /* no. objects */
51             int numClusters, /* no. clusters */

```

```

52     double threshold,      /* minimum fraction of objects that change membership */
53     long loop_threshold, /* maximum number of iterations */
54     int *membership,      /* out: [numObjs] */
55     double *clusters)    /* out: [numClusters][numCoords] */
56 {
57     int i, j, k;
58     int index, loop = 0;
59     double timing = 0;
60
61     double delta;          // fraction of objects whose clusters change in each loop
62     int *newClusterSize; // [numClusters]: no. objects assigned in each new cluster
63     double *newClusters; // [numClusters][numCoords]
64     int nthreads;         // no. threads
65
66     nthreads = omp_get_max_threads();
67     printf("OpenMP Kmeans - Reduction\t(number of threads: %d)\n", nthreads);
68
69     // initialize membership
70     for (i = 0; i < numObjs; i++)
71         membership[i] = -1;
72
73     // initialize newClusterSize and newClusters to all 0
74     newClusterSize = (typeof(newClusterSize))calloc(numClusters, sizeof(*newClusterSize));
75     newClusters = (typeof(newClusters))calloc(numClusters * numCoords,
76         sizeof(*newClusters));
76
77     // Each thread calculates new centers using a private space. After that, thread 0 does
78     // an array reduction on them.
78     int *local_newClusterSize[nthreads]; // [nthreads][numClusters]
79     double *local_newClusters[nthreads]; // [nthreads][numClusters][numCoords]
80
81     /*
82      * Hint for false-sharing
83      * This is noticed when numCoords is low (and neighboring local_newClusters exist
84      * close to each other).
85      * Allocate local cluster data with a "first-touch" policy.
86      */
86     // Initialize local (per-thread) arrays (and later collect result on global arrays)
87     for (k = 0; k < nthreads; k++)
88     {
89         local_newClusterSize[k] = (typeof(*local_newClusterSize))calloc(numClusters,
90             sizeof(**local_newClusterSize));
90         local_newClusters[k] = (typeof(*local_newClusters))calloc(numClusters * numCoords,
91             sizeof(**local_newClusters));
91     }
92
93     timing = wtime();
94     do
95     {
96         // before each loop, set cluster data to 0
97         for (i = 0; i < numClusters; i++)
98         {
99             for (j = 0; j < numCoords; j++)
100                 newClusters[i * numCoords + j] = 0.0;

```

```

101         newClusterSize[i] = 0;
102     }
103
104     // reset delta before each iteration; it will be updated via reduction in the
105     // parallel region
106     delta = 0.0;
107
108     /*
109      * TODO: Initialize local cluster data to zero (separate for each thread)
110      *
111      * We now use an OpenMP parallel region where:
112      * - Each thread zeroes its own local_newClusterSize/local_newClusters.
113      * - The object loop is distributed with 'omp for' and 'reduction(+ : delta)'.
114      * - A single thread reduces the per-thread local arrays into the shared arrays.
115      */
116 #pragma omp parallel private(i, j, k, index)
117 {
118     int tid = omp_get_thread_num();
119     int T   = omp_get_num_threads(); // actual number of threads in this team
120
121     /* per-thread zeroing (first-touch initialization of local cluster data) */
122     for (i = 0; i < numClusters; i++)
123         local_newClusterSize[tid][i] = 0;
124     for (i = 0; i < numClusters * numCoords; i++)
125         local_newClusters[tid][i] = 0.0;
126
127     // Distribute objects across threads and compute per-thread contributions.
128     // delta is accumulated using a reduction to avoid atomics on a shared
129     // variable.
130 #pragma omp for reduction(+ : delta)
131     for (i = 0; i < numObjs; i++)
132     {
133         // find the array index of nearest cluster center
134         index = find_nearest_cluster(numClusters, numCoords,
135                                     &objects[i * numCoords], clusters);
136
137         // if membership changes, increase delta by 1
138         if (membership[i] != index)
139             delta += 1.0;
140
141         // assign the membership to object i
142         membership[i] = index;
143
144         // update new cluster centers : sum of all objects located within (average
145         // will be performed later)
146         /*
147          * TODO: Collect cluster data in local arrays (local to each thread)
148          * Replace global arrays with local per-thread
149          */
150         local_newClusterSize[tid][index]++;
151         for (j = 0; j < numCoords; j++)
152             local_newClusters[tid][index * numCoords + j] += objects[i * numCoords
153 + j];
154     }

```

```

151
152     /*
153      * TODO: Reduction of cluster data from local arrays to shared.
154      * This operation will be performed by one thread
155      *
156      * Here we use 'omp single' so that exactly one thread accumulates
157      * all per-thread local arrays into the shared newClusterSize/newClusters.
158      */
159 #pragma omp single
160 {
161     for (k = 0; k < T; k++) // only sum over the threads actually in this
team
162     {
163         int *srcS = local_newClusterSize[k];
164         double *srcC = local_newClusters[k];
165         if (!srcS || !srcC)
166             continue;
167         for (i = 0; i < numClusters; i++)
168         {
169             newClusterSize[i] += srcS[i];
170             for (j = 0; j < numCoords; j++)
171                 newClusters[i * numCoords + j] += srcC[i * numCoords + j];
172         }
173     }
174     /* implicit barrier after single */
175 } /* end parallel region */

176
177 // average the sum and replace old cluster centers with newClusters
178 for (i = 0; i < numClusters; i++)
179 {
180     if (newClusterSize[i] > 0)
181     {
182         for (j = 0; j < numCoords; j++)
183         {
184             clusters[i * numCoords + j] = newClusters[i * numCoords + j] /
newClusterSize[i];
185         }
186     }
187 }

188
189 // Get fraction of objects whose membership changed during this loop. This is used
as a convergence criterion.
190     delta /= numObjs;

191
192     loop++;
193     printf("\r\tcompleted loop %d", loop);
194     fflush(stdout);
195 } while (delta > threshold && loop < loop_threshold);
196     timing = wtime() - timing;
197     printf("\n nloops = %3d (total = %7.4fs) (per loop = %7.4fs)\n", loop, timing, timing
/ loop);

198
199     for (k = 0; k < nthreads; k++)
200     {

```

```
201     free(local_newClusterSize[k]);
202     free(local_newClusters[k]);
203 }
204 free(newClusters);
205 free(newClusterSize);
206 }
207
208 }
```